

The Governed Medi-Cal Enterprise

From MITA maturity to governed delegation: holding every material Medi-Cal decision at a known, evidenced, revocable level of machine authority — and proving it.

Medi-Cal is about to be asked to do more determining, verifying and reconciling than at any time since the Affordable Care Act expansion — and to do it in the fiscal year its principal financing mechanism is dismantled. That is not an argument for automation. Automation is the easy half, and it is the half that has already failed in five other jurisdictions.

The hard half is **delegation**: deciding which decisions a machine may make, on what evidence, under whose authority, and by what mechanism that authority is taken back. California has no instrument for that. Neither does MITA — and CMS knows it, which is why the maturity model is being rebuilt around outcomes and metrics while its successor page still reads *Under Construction*.

This paper proposes the instrument: four rungs of machine authority, gated by four tiers of evidence, applied to one business process and one action class at a time. It says plainly what can never climb.

WHAT THIS PAPER DOES NOT CLAIM

Not "substantially autonomous in five years." The gate arithmetic in §5 refutes it. The achievable state is A2-dominant with a narrow A3 tail.

Not "already funded." 90/10 and 75/25 are match rates, not appropriations. The state share is real money. §10 publishes it.

Not a running system. The reference implementation cited in §5 has never processed a Medi-Cal transaction. Its live-integration count is zero.




FIGURE 1 The Medi-Cal North Star. One state-owned evidence record beneath everything; a controller that conducts the business processes above it and writes every action and its evidence back down; a registry in which each agent declares its scope, its data tier, its autonomy rung and its accountable owner; and governance gates that sit *in* the decision loop rather than after it.

§1 • THE VISION

Eight things that become true

Each is an operating outcome, not a technology. None of them is reached by buying a platform, and none of them is reached at once.

1 • Standing eligibility

Medi-Cal enrollment becomes a status the enterprise continuously evidences on the beneficiary's behalf, not a claim they must periodically re-prove to keep coverage they never lost. From a renewal packet mailed on a calendar and worked by a county eligibility worker — with coverage lost for procedural rather than substantive reasons — to ex parte renewal as the default, and a human engaged where the evidence conflicts or the judgment is genuinely required.

● EE01–EE04 · ME01–ME02

2 • The three-way match

Authorization, encounter and payment are held in continuous agreement as one reconciled thread, so what was authorized, what was delivered and what was paid are never three separate accounts that only a retrospective audit can reconcile. This is the anchor, and §7 is its deep dive.

● OM07 · OM29 · OM14 · OM27 · FM

3 • Physician-final, machine-fast

The evidence a medical-necessity decision requires is assembled and checked the moment the request arrives, so a clinician's time is spent only on requests that actually turn on clinical judgment — and every denial remains a clinician's act. SB 1120 is not a limitation this paper concedes. It is the constraint that makes the element credible.

● CM07 · CM08 · CM09

4 • Credential once, verified always

A provider or community-based organization proves itself once and is thereafter continuously verified against authoritative sources, rather than re-proving what the state already holds at every revalidation. For a small Community Supports provider, administrative burden is a participation decision, not an inconvenience.

● PM01–PM03 · EE05–EE08

5 • No one lost at the seam

The transitions where Medi-Cal loses people — release from custody, hospital discharge, plan transition, county move, the dual-eligible boundary — are watched continuously and closed by the enterprise, rather than discovered by whichever care manager happens to notice.

● CM01 · CM02 · CM06 · ME03

6 • Integrity before payment

Improper payment is surfaced while the pattern is forming rather than recovered years later from providers who may no longer exist. Recoveries return the federal share to CMS. This is an avoided-cost and error-rate outcome, and never a funding source — a distinction the Legislative Analyst will test first.

● PE01–PE05 · FM recovery

7 • The state owns the proof

DHCS can evidence, contemporaneously and on demand, what it determined, authorized, substantiated and paid — and on what basis — independent of whichever vendor or application is running in a given budget cycle. Stated as "a data platform," this is a procurement. Stated as "the state can prove what it did," it is audit defense, federal-match integrity and renewal defensibility.

● OM28 · cross-cutting

8 • Earned, and revocable

No Medi-Cal business process operates at a higher rung of machine authority than its own accumulated evidence record supports, and every rung can be walked back on evidence — quickly, by one named person, without a committee.

● PL · PE · cross-cutting

THE ONE THAT IS DIFFERENT

Seven of these are outcomes California would want in any decade. The eighth is the one that makes the other seven safe, and it is the only one that does not exist anywhere in American Medicaid today. **A state that automates without it has not modernized; it has simply moved its judgment somewhere it cannot see.**

§2 · THE CLOCK

Five obligations, one day, no new money

The case for acting now is not that the technology arrived. It is that the workload arrives on a date certain, and the revenue leaves on the same date.




FIGURE 2 The obligations that dominate 2026–2032. Horizon extends to 2034; nothing material is scheduled after 2029. Every date shown is confirmed against the Federal Register, the eCFR, DHCS publications or the enacted Budget Act.

THE CONVERGENCE

On one day — 1 January 2027 — DHCS must have running community-engagement verification for the adult group, six-month redeterminations for that same group, shortened retroactive-coverage windows, address verification from reliable data sources, and quarterly Death Master File matching. On that same day the four CMS-0057-F APIs go live for fee-for-service and for every Medi-Cal managed care, mental health and drug Medi-Cal plan.

There is no version of that date California staffs its way through. And the funder has already written the preference into its own guidance: a state **must first attempt to use reliable information available to the state** before asking a human being for a document. That sentence is an automation mandate in eligibility, issued by CMS.

THE PINCER

The MCO tax expires 31 December 2026. The Legislative Analyst projects roughly \$10 billion in annual operating deficits through 2029–30. The enacted 2026–27 Budget Act books a conforming replacement tax rising to \$2.3 billion; the Analyst's own projection for the same instrument was in the mid-tens of millions. That spread — roughly \$2 billion a year — rests on a federal approval this paper cannot verify, and it should be read as a range, not a number.

THE MEASUREMENT REGIME

From 2028 CMS stops accepting attestations about California and starts testing them. Wait-time standards bite on the 2028 rating period. Secret-shopper surveys validate the network data on the 2029 one. The Quality Rating System publishes by December 2028. And from October 2029 an eligibility error rate above three percent triggers disallowance, with the good-faith waiver restricted.

Read together, these are not four compliance items. They are an **evidence regime**. By 2029 the only claim California will be permitted to make about its own operations is one it can substantiate per process, with a measured error rate and a confidence interval. That is precisely the currency the ladder in §5 is denominated in — which is why building the instrument is not an addition to the compliance workload. It is the compliance workload, done once.

AND ONE THING THAT IS NOT SCHEDULED

DHCS has published no dated modular roadmap for the Medi-Cal Enterprise System. The behavioral-health consolidation, MEDS modernization and PAVE all appear without go-live dates. **The absence is the argument.** A per-process evidence record is the cheapest thing that would fill it.

§3 · THE CEILING

What Medi-Cal may never automate

Every serious proposal of this kind begins with what becomes possible. This one begins with what does not, because the limits are the reason to believe the rest.

An **adverse determination affecting a beneficiary's eligibility, coverage, or payment may never operate above human-in-the-loop**. No volume of evidence promotes it, and no fiscal pressure justifies it. The rule is evaluated *before* the autonomy rung is read, and it overrides it — not as a policy preference the state could revisit in a harder budget year, but as the intersection of three independent legal constraints, each of which forecloses the question on its own.

CONSTITUTIONAL DUE PROCESS

Goldberg v. Kelly is not merely echoed in Medicaid; it is incorporated by name — **42 CFR § 431.205(d)** requires the hearing system to meet that decision's due-process standards. Three mechanisms follow, each independently fatal to machine execution. **§ 431.210** requires "a clear statement of the specific reasons supporting the intended action": an adverse action whose actual causal basis cannot be stated as a reason to the beneficiary is not noticeable, and therefore not lawfully takeable. **§ 431.211** requires notice at least ten days before the date of action, which creates an irreducible human step. And **§ 431.240(a)** requires hearings by officials "who have not been directly involved in the initial determination" — which presupposes an identifiable initial determiner. **Autonomy that erases the identity of the decider also erases the conflict check the regulation depends on.**

MANAGED CARE CARRIES THE SAME STRUCTURE

The cleanest federal text of all is **§ 438.210(b)(3)**: a decision to deny a service authorization, or to authorize less than requested, must be made **by an individual** with appropriate expertise. Not by a system configured by such an individual. By an individual. **§ 438.400(b)** supplies the definition and **§ 438.406(b)(2)** keeps the appeal separate from the decision it reviews.

CALIFORNIA HAS ALREADY LEGISLATED THE ANSWER

SB 1120 (Stats. 2024, ch. 879, effective 1 January 2025) provides that an AI or algorithmic tool "shall not deny, delay, or modify health care services based, in whole or in part, on medical necessity," and that such a determination "shall be made only by a licensed physician or a licensed health care professional competent to evaluate the specific clinical issues involved."

Read it precisely, because the precision is the gift. SB 1120 does not ban AI in utilization management. **It bans AI from being the basis of an adverse outcome, and says nothing about favorable ones.** That asymmetry by valence is the axis this paper's ladder is built on, and it is the legal reason approvals may climb while denials may not. "In whole or in part" is severe enough that for a medical-necessity denial the safe posture is stronger than human-in-the-loop: the machine assembles the evidence and does not propose the disposition at all. This paper names that rung **A1-restricted**.

§ 1367.01 reaches Medi-Cal managed care plans as Knox-Keene licensees, but not DHCS's own fee-for-service review — so **DHCS should adopt the standard for its own operations as policy where the statute does not compel it.**

The Governed Medi-Cal Enterprise

THE CEILING TEST

Evaluated *before* the autonomy rung is read. An action is ceiling-bound — maximum A1, or A1-restricted where SB 1120 applies — if **any** of the following is true:

1. It is an adverse benefit determination within 42 CFR § 438.400(b).
2. It denies, terminates, reduces or suspends eligibility, or denies an application, under 42 CFR 431 Subpart E.
3. It creates or increases beneficiary financial liability.
4. It denies, delays or modifies a service based in whole or in part on medical necessity.
5. **It shifts the burden of production to the beneficiary**, or records a non-confirming result in an automated verification — an ex parte failure, a data-match discrepancy, a deadline non-response, or any action that starts a discontinuance clock.
6. It is functionally equivalent to any of the above by any other name.

WHY LIMB FIVE EXISTS

A ceiling drawn around the *signature* on an adverse action is a story about a ceiling. Ex parte failure, non-confirmation, and a request for information with a deadline are all favorable-or-neutral by label — and every one of them is where people are actually lost. Limb five is why this is an adverse-*effect* gate, not an adverse-*action* gate.

ANTI-CIRCUMVENTION

Pend-as-denial. A pend that produces the effect of a denial is a denial; auto-pend needs a hard cap and may never auto-resolve adversely on timeout.

Partial approval. Authorizing less than requested is adverse.

§4 · BUSINESS ARCHITECTURE

MITA's ladder, and the one CMS is already building

MITA gave state Medicaid a common language — ten business areas, eighty business processes, and a maturity model rating each of them on six qualities. Read the qualities and notice what is absent. Timeliness, data accuracy, effort, cost effectiveness, accuracy of results, value to stakeholders: **every one is a property of the output. None is a property of the decider.** There is no quality for who or what made the decision, none for whether it can be reconstructed, and none for how much human judgment the process still consumes.

The ladder itself confirms it. ML4 is about intrastate information exchange and quantitative measurement. ML5 speaks of "national (and international) interoperability improvements that maximize automation of routine operations." That is a real automation clause and this paper does not pretend otherwise — but its object is *routine operations*, the rule-determinate residue, and its organizing telos is interoperability. Nowhere across ML1 to ML5 is there a construct for who holds a decision, what evidence justifies moving that authority, or how it is taken back.

That gap is not a problem for California. It is an opening — because CMS is rebuilding the model right now, and has said what it will be made of.

MITA is not being retired: under **42 CFR § 433.112(b)(11)**, advancing in MITA maturity remains a condition of the 90 percent match, and CMS reaffirms that the State Self-Assessment is still required. But the Governance Board has been rebuilding the framework since 2023, MITA 4.0 has no release date, and the maturity-model page in CMS's own MITA 4.0 repository still reads *Under Construction*. Meanwhile CMS's live operational guidance — Streamlined Modular Certification, and the MES template set mandatory from July 2026 — has moved to outcomes and metrics and no longer cites MITA at all.

CMS has told states what the successor is built from: a maturity model "tied to outcomes and metrics," with each capability defined by outcomes, roles, business processes, information and technology. **A governed autonomy ladder, evidenced per business process, is the most direct expression of that specification anyone has proposed.** California can arrive at MITA 4.0 with the instrument already running.

THE CAPABILITY MODEL · 31 BUSINESS CAPABILITIES IN SIX LAYERS

The paper's own construct, not a MITA artifact. Layers are ordered by dependency, and each capability is something the enterprise can *do*, never a system it can buy. The boundary rule: **E holds facts and policy · O moves work · D exercises judgment · G touches people · A watches D and O · F determines and moves money.**

E · EVIDENCE & MEMORY

- E1 Entity Resolution
- E2 Longitudinal Record Assembly
- E3 Evidence Ledger
- E4 Policy-as-Code Registry
- E5 Decision Provenance & Replay
- E6 Consent & Minimum-Necessary
- E7 Disclosure Resistance

O · ORCHESTRATION

- O1 Case & Work Orchestration
- O2 Handoff Contract Enforcement
- O3 Exception & Escalation
- O4 Autonomy Tier Control Plane
- O5 Human Attention Allocation

D · DECISIONING

- D1 Eligibility & Benefit Determination
- D2 Medical Necessity & Authorization
- D3 Claim & Encounter Adjudication
- D4 Anomaly & Integrity Scoring
- D5 Adverse-Effect Gate
- D6 Policy Change Simulation

G · ENGAGEMENT

- G1 Member Interaction
- G2 Provider, Plan & County Interaction
- G3 Outreach & Deadline Rescue
- G4 Notice, Language & Accessibility

A · ASSURANCE & GOVERNANCE

- A1 Autonomy Evidence Record
- A2 Continuous Control Monitoring
- A3 Appeals & Fair Hearing Support
- A4 Disparate-Impact Surveillance
- A5 Change Control & Rollback
- A6 Oversight & Audit Interface

F · FINANCIAL

- F1 Rate & Capitation Administration
- F2 Payment Execution & Settlement
- F3 Recovery & Cost Avoidance
- F4 Federal Claiming & Match Defense

WHERE THE VALUE ACTUALLY LEAKS

Six end-to-end value chains cross these business areas: a person becomes covered and stays covered · a provider becomes trusted · a service is authorized and delivered · **a claim or encounter becomes a settled, defensible payment** · an integrity signal becomes a prevented loss · a policy change becomes operating reality. **Value leaks at the handoffs** — precisely the unit of analysis MITA's ladder does not have

READ THIS AS DESIGN, NOT STANDARD

The capability model, the rungs, the tiers and every threshold here are the authors' construct, carrying no external authority — stated precisely so they can be argued with, not mistaken for

§5 · THE TWIN LADDER

Authority on one axis, evidence on the other

A business process may never operate above the autonomy rung that its data tier supports. That single rule converts "are we confident enough to automate this?" — an opinion — into "has this process's evidence reached the tier that permits the next rung?" — a fact that can be measured, published, and refused.

FOUR RUNGS OF MACHINE AUTHORITY

	MACHINE PRODUCES	ON HUMAN SILENCE	HUMAN'S UNIT OF WORK	IN-BAND SAFETY
A0 Assist	an artifact	nothing	the decision	an independent human path
A1 HITL	a proposal	nothing	the instance	suspension
A2 HOTL	proposal + timer	the action	the population	window, stop, sampling
A3 Auto	the effect	the action	the policy	none

The distinguishing test between A1 and A2 is the default on silence: at A1 silence means nothing happens; at A2 silence means the machine's action happens. Between A2 and A3 it is whether a stop window exists at all. **A3 has no in-band safety whatsoever** — which is why an action class may reach it only where the worst realistic failure is a recoverable error. If you cannot describe the compensating transaction, the class cannot be A3.

FOUR TIERS OF EVIDENCE

	THE QUESTION THE DATA CAN ANSWER
D0 Fragmented	What does system X say?
D1 Conformed	What is true about this entity — who said so, and when did we learn it?
D2 Reconciled & contestable	What is true when the parties disagree, how wrong are we, and who has told us we are wrong?
D3 Settlement-grade	What was true at the moment of decision, provably, in a package an adjudicator will accept?

A tier is not an adjective about data quality. It is a claim about what interrogation the evidence survives. Three rules make it auditable: a process's tier is the **minimum** over its decision-critical inputs, never the average; tier is attained per (process, action class, scope), so enterprise-wide tier claims are meaningless by construction; and **tier is not monotone in time** — a schema change, a new external match, or a policy change that makes a dormant field decision-critical can drop a tier, and must therefore drop the rung.

The twin ladder: data tier gates autonomy rung




FIGURE 3 A process may occupy any cell on or above the frontier. Tracks show 2026 → 2029 → 2034; tier placements are design judgements from the current-state research, *not measurements* — replacing them with measured values is the whole purpose of the instrument in §9. Adverse determinations within processes 3 and 5 remain at A1 permanently under 42 CFR 431 Subpart E and SB 1120 regardless of tier attained; only the favourable and neutral action classes are plotted. Note process 6: the behavioral health record does not reach D2 by 2034, and the reasons are fiscal and legal rather than technical.

§6 · THE RECORD

Every catastrophe was an evidence failure, not a model failure

Across five jurisdictions and two decades, the documented disasters of public-benefits automation share one structure — and in none of them was the automation sophisticated.

CASE	WHAT THE AUTOMATION DID	THE EVIDENCE READING
Michigan MIDAS 2013–2015 · <i>Bauserman</i>	Auto-adjudicated unemployment fraud from mismatches between claimant-reported earnings and employer wage records, flagging any discrepancy however trivial. Notice reached claimants only through a portal they had no reason to open; non-response defaulted to an adverse determination . The Auditor General reviewed 22,000 determinations and found 93 percent involved no fraud.	Two parties asserting the same fact at different temporal granularities, with no reconciliation layer, no versioned precedence rule, and a contest channel that structurally did not reach the affected person. D1 at best, run at A2 — silence executing as consent.
Australia, Robodebt 2015–2019 · <i>Prygodicz</i> (No 2)	Averaged an annual tax-office income figure across twenty-six fortnights, tested it against fortnightly reported income, raised a debt on the difference — and reversed the onus onto the recipient. The Federal Court found no proper legal basis; roughly 648,000 group members, about A\$1.763 billion in debts withdrawn.	Averaging is a <i>conformance</i> operation — so this is a D1 failure, not a D0 one. The enterprise successfully assembled a record and, in assembling it, fabricated a fact no source had asserted . The debts had to be withdrawn wholesale because the Commonwealth could not tell the correct ones from the incorrect: the price of missing restatement lineage.
Idaho <i>K.W. v. Armstrong</i> , 2016	An Excel budget tool set annual Medicaid service funding for roughly 4,000 people with developmental disabilities. The formula was fitted on 733 of 3,512 participant records — the other 79 percent discarded as unreliable. Never recalibrated, never audited. The court held it unconstitutionally arbitrary.	The purest tier failure on record. No model improvement rescues a formula fitted on a corrupted remnant of its own population — the case to put in front of anyone proposing to solve an evidence problem with a better algorithm.
Arkansas & Indiana <i>Ledgerwood</i> 2017 · <i>Perdue</i> 2012	An algorithmic tool replaced nurse discretion in setting attendant-care hours, producing large reductions; the courts held the state must explain the specific factors that produced a cut. Indiana's automated eligibility processing issued denials carrying only generic reason codes.	Not data quality but provenance and contestability : neither state could say, in a notice, which inputs produced the outcome. The machine held the disposition but not <i>the reason as a data structure</i> . The same requirement binds DHCS under § 431.210 — a rationale that cannot become a notice is not a rationale .
Tennessee TEDS <i>A.M.C. v. Smith</i> , 2024	The system failed for months to load Social Security benefits, disability status and marital status, so it could not evaluate disability-related coverage categories; during migration, enrollees were sorted into wrong households or dropped entirely.	D0 masquerading as D1 . An assembled record, silently incomplete on exactly the fields deciding alternate-basis eligibility, with nothing detecting the incompleteness. This is <i>ex parte</i> redetermination — a D2 operation — run on data that had not attained D1 coverage. It is the same process California must run at six-month cadence from January 2027 .

THE DIAGONAL: A2 RUNNING ON D1 EVIDENCE

The forbidden region is a staircase, not a gradient — a process at D1 attempting A2 is not "mostly safe," it is running the configuration above. **Its circuit breaker cannot be calibrated**: the breaker trips when the executed-action distribution moves beyond a pre-registered limit, which needs a stable baseline and a known error rate so the system can tell *the world changed* from *our data changed*. On D1 evidence a schema change or a mass restatement moves that distribution exactly as a real behavioral shift would — leaving limits so wide the breaker never trips, or so narrow it is switched off within a quarter. Since A2's failure mode is precisely that nobody is looking, a decorative breaker means no safety at all.

Silence becomes uninterpretable, and A2 runs on silence. **Correlated errors defeat sampling**, because a tier defect strikes every record sharing the defective attribute while the reviewer checks dispositions, not the data beneath them. And **the compensating transaction works on one case and fails on the cohort**.

The Governed Medi-Cal Enterprise

THE GENERALISATION

Temporal misalignment between two sources describing one fact. An unusable population under a fitted rule. Provenance too thin for a lawful notice. Silently incomplete assembly presented as complete.

Every one is an evidence-tier failure. The tier rule is the cheapest safety control available: it forbids, in advance and on inspectable evidence, the exact configuration that produced all of them.

§7 · THE ANCHOR

A claim or encounter becomes a settled, defensible payment

One value chain, two disciplines, six handoffs. Each discipline discovers the other's errors late — and the losses do not sit inside the decisions. They sit between them.

FEE-FOR-SERVICE · CA-MMIS



MANAGED CARE · ENCOUNTERS



FIGURE 4 Two disciplines that only meet in a rate cell. Numbered markers are the handoffs where value leaks, described opposite. The shaded boxes hold the enterprise's largest concentration of human judgment and its least-instrumented one.

WHERE THE VALUE LEAKS

- 1 · Provider → claim.** Billing error is manufactured before the enterprise ever sees the transaction, and the only feedback instrument is a denial arriving weeks later.
- 2 · Edit → suspense.** The suspended-claim queue is the largest concentration of human judgment in the program and its least-instrumented: cycle time, resolution consistency and rework rate are not first-class managed quantities.
- 3 · Payment → federal claiming.** Reconstructing *why* a specific dollar was claimed at a specific match rate, years later, is a project rather than a query. That is the deferred cost of an enterprise with no evidence ledger.
- 4 · Plan adjudication → encounter.** The plan pays on its own record and *then* reports, so the encounter is a derived artifact rather than the transaction — and the disagreement between what the plan adjudicated and what the state accepted is visible only in aggregate.
- 5 · Rejection → correction.** Rejected encounters loop back on a batch cadence. That loop's latency is the latency of the state's own view of what care was delivered.
- 6 · Encounters → rate setting.** Incomplete or late encounters propagate into capitation rates. This is the most expensive leak in the chain, because the error compounds forward into every future payment.

WHY THIS IS THE RIGHT PLACE TO LEARN

Not because it is where the money is — because it is where the *evidence* is. Every claim is a labelled decision carrying a disposition, a remittance advice, and, decisively, **a dispute channel in which the disputing party is motivated by money.** That is the strongest ground-truth signal anywhere in the enterprise, and it accumulates faster here than anywhere else.

So the claim is narrow and worth stating precisely: **clean-line adjudication is where the method is proven, not where the value is.** A paper that anchored on value would anchor on eligibility — and would be proposing to learn governance on the population least able to absorb a mistake.

AND ONE THING THIS PAPER WILL NOT SAY

It is true in law that a provider payment denial is not a beneficiary due-process event. It would be catastrophic to build on it. For a health centre a denied encounter is absorbed whole; for county behavioral health, a program that gets faster and more consistent at *not paying* is a solvency event. So the rule runs the other way: **no action class that denies or reduces a payment to a provider climbs above A1**, and the provider remedy chain — electronic appeal, a committed resolution date on suspended claims, a documentation loop that ends timeout denials — climbs first.

§8 · THE CLIMB

One decision, four rungs, eight years

Take the highest-volume decision in the program — the payment disposition on a single line of a fee-for-service professional claim — and walk it up. The ceiling analysis comes first, because that is the order the instrument runs in.

DISPOSITION	CEILING-BOUND?	MAX RUNG	BASIS
Pay as billed — clean line	No	A3 eligible	Favorable or neutral; no adverse determination; reversible by recoupment within the cycle
Pay adjusted — pricing or edit	Only if no beneficiary liability results	A2	Provider-facing adjustment; verify no cost-sharing effect. A reduction to provider payment is separately floored at A1 by policy
Pend for information	No — with a hard cap	A2	Auto-pend permitted only with a maximum cumulative duration; on expiry it routes to a human and may <i>never</i> auto-resolve adversely on timeout
Deny — adverse benefit determination	Yes	A1 permanently	42 CFR § 438.400(b); 42 CFR 431 Subpart E where eligibility-linked
Deny on medical necessity	Yes	A1-restricted	SB 1120 — the tool may not deny, delay or modify based in whole or in part on medical necessity, so it may not even propose the disposition

YEARS 1–2 · A0, AND THE YEAR THAT PRODUCES NO AUTONOMY

For each suspended line the machine assembles the adjudication packet — eligibility on date of service, benefit and authorization status, provider enrollment, pricing reference, duplicate history, edit rationale — and produces a ranked candidate disposition with a stated reason. It writes nothing to the claims system, and the examiner may ignore it entirely.

What matters is the instrumentation, not the assistance: **every machine candidate is paired with the examiner's actual disposition and written to the evidence ledger.** A0 without recorded pairs can never earn A1, because the concordance evidence cannot be reconstructed afterwards.

And here the least comfortable admission belongs. The baselines the gate depends on — suspense disposition mix, appeal overturn, human-to-human concordance on the same line — **are not published in California, and in at least one case have never been measured.** No action class enters the ladder until its own baseline is. **Year one delivers a number, not a capability** — which is also the only Year One deliverable in this field the Legislative Analyst can check.

YEARS 3–4 · A1

The machine emits a complete, executable proposed disposition and suspends; the examiner's work flips from production to review. Silence produces nothing. And reviewers whose modification rate is statistically indistinguishable from zero have their decisions **excluded from the promotion evidence pool**, with the exclusion count printed in the package — because without that control, "the human agreed" degenerates into "the human clicked."

YEAR 5 · THE FORK, AND THE REFUSAL

Two promotion packages are prepared. Only one can succeed.

Clean-line pay-as-billed → A2. Granted on evidence: sustained concordance with a confidence bound, calibration with no over-confident decile, no subpopulation disparity outside the pre-registered band holding case-mix constant, an overturn rate non-inferior to the human baseline, and a reversibility attestation naming a recoupment path that has actually executed.

Denials → A2. Denied, and denied permanently. Not for want of evidence. The ceiling test fires before the tier or the rung is read. A denial constituting an adverse benefit determination stays at A1 at ninety-nine percent concordance, and stays at A1 at one hundred.

There is also a technical precondition that must not be waved through. A2's entire safety rests on an intervention timer firing. **No action class may be promoted to A2 until the durable execution substrate is running in production and its timer-survival property has been verified under restart.** A timer that dies on deploy converts A2 into A3 silently, and nobody finds out until the quarter closes.

YEARS 6–8 · A2, AND A NARROW A3

Cycle time decouples from headcount for the qualified clean line, and the throughput ceiling stops being reviewer capacity. Every other disposition — every denial, every partial, every medical-necessity question — continues to route to A1 in the same process and the same transaction.

That is the honest picture of Medi-Cal in 2034: **A2-dominant, with an A3 tail, and a permanent floor under every decision that defines the program's relationship to a person.** It is a smaller claim than the one usually made in this field. It is also the only one this paper's own gate arithmetic supports.

§9 · THE INSTRUMENT, THE MONEY, THE NEXT TWELVE MONTHS

What California would actually do

THE MEASUREMENT INSTRUMENT

One page per (process, action class, scope) — the same unit the ladder tiers — carrying twelve indicators, three per tier: false-merge and false-split rates reported separately and never netted; lineage completeness; source-to-record latency per feed; conflicts detected and closed under a versioned precedence rule; field-level error with a confidence interval; challenges per thousand decisions and the proportion upheld; chain-verification date; replay fidelity; and appeal-package assembly time from an unannounced drill.

Nine of the twelve are views over data DHCS already produces — payment error measurement, encounter quality measurement, the prior-authorization metrics already published under CMS-0057-F, grievance and appeal reporting, and interface SLAs. The instrument is a lens, not a collection, and it rides an operational reporting vehicle that is already mandatory. **One indicator — replay fidelity — is genuinely new and costs real money.** It should be justified on appeal defence and audit economics rather than dressed in a mandate that does not carry it.

One interpretive rule travels with the contestability indicator, and it reverses the sign on the metric that failed in Australia. **A falling challenge rate is a warning, not an achievement.** Automation that produces less comprehensible notices, or moves faster than people can respond, lowers the appeal rate — and a gate rewarding that would promote a process for silencing its own critics. So appeal overturn is demoted to monitoring and replaced, as gate evidence, by independent shadow adjudication on a sample that includes decisions nobody appealed.

THE MONEY, AS A BUDGET ANALYST WOULD STATE IT

Enhanced federal match is real and intact: 90 percent for design and development of an approved system, 75 percent for its operation, both unchanged in the current eCFR, with modularity, reuse and no proprietary lock-in as conditions — and MITA advancement itself a condition of the ninety. But **90/10 and 75/25 are match rates, not appropriations.** California must appropriate the state share first; the 75 applies only after certification; and the governance apparatus this paper spends most of its words on — promotion authority, sampling, shadow adjudication, drills, oversight — is program administration at **50/50.**

The honest framing is *already at ninety cents on the federal dollar, and already inside an advance planning document California must file anyway.* It is not free — and the state share competes directly with the enacted reductions to Enhanced Care Management and Community Supports, which should be said in the first meeting rather than discovered in the third.

And the return should be described precisely. **This program does not produce General Fund savings inside its horizon, because DHCS does not hold the labor.** The work sits with counties, with the fiscal intermediary, and inside capitation. What it produces is avoided disenrollment, avoided disallowance once the three-percent error threshold binds in 2029, and avoided cutover failure. Anyone promising headcount is promising what the org chart cannot deliver.

1 · PUT IT IN THE CONTRACT

The fiscal intermediary contract determines whether California can ever see its own claims decisions. Per-decision provenance emission as a deliverable on a state-owned schema; state ownership of suspense-queue telemetry as a reported performance measure; a contractual right to place a state-owned ceiling gate in the disposition path; portability of the evidence record on exit. **The requirement costs almost nothing in a solicitation and almost everything as a change order** — and no rung above A0 is reachable in claims processing until it lands.

2 · PUBLISH THE REGISTER

Every action class with its data tier, its rung, its evidence volume, its concordance estimate and interval, its rubber-stamp exclusion count, its demotion events and its next gate date — quarterly to the Legislature, with a four-line portfolio scorecard above it. The number that proves the instrument is real is not the count of promotions granted. **It is the count refused,** and publishing it costs nothing.

3 · BARGAIN BEFORE YOU BUILD

No promotion touching a represented classification proceeds without completed meet-and-confer. A document carrying thirty-one capabilities and an eight-year roadmap that never mentions the workforce performing every A0 and A1 step is a labor program in a governance costume. Make the reinvestment promise falsifiable: publish a **human attention floor** in measured minutes per ceiling-bound case, pre-registered before any promotion, with automatic demotion if measured review time falls beneath it.

California will be measured on the accuracy of its own evidence from 2029, on a schedule it does not control. The only open question is whether it measures itself first.